

Import Libraries

```
In [1]: import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
%matplotlib inline
```

Upload dataset

```
In [2]: data = pd.read_csv('C:\\Users\\sundara.rao.ext\\Desktop\\SUNDAR\\Data Science
\\DL & NLP\\Machine Learning\\Machine Learning\\Credit_Card_Default_Classifica
tion\\creditcard.csv')
```

Analysis the data set

```
In [3]: # View the shape of dataset
data.shape
```

```
Out[3]: (284807, 31)
```

```
In [4]: # View the type of data
data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 284807 entries, 0 to 284806
Data columns (total 31 columns):
Time          284807 non-null float64
V1            284807 non-null float64
V2            284807 non-null float64
V3            284807 non-null float64
V4            284807 non-null float64
V5            284807 non-null float64
V6            284807 non-null float64
V7            284807 non-null float64
V8            284807 non-null float64
V9            284807 non-null float64
V10           284807 non-null float64
V11           284807 non-null float64
V12           284807 non-null float64
V13           284807 non-null float64
V14           284807 non-null float64
V15           284807 non-null float64
V16           284807 non-null float64
V17           284807 non-null float64
V18           284807 non-null float64
V19           284807 non-null float64
V20           284807 non-null float64
V21           284807 non-null float64
V22           284807 non-null float64
V23           284807 non-null float64
V24           284807 non-null float64
V25           284807 non-null float64
V26           284807 non-null float64
V27           284807 non-null float64
V28           284807 non-null float64
Amount        284807 non-null float64
Class         284807 non-null int64
dtypes: float64(30), int64(1)
memory usage: 67.4 MB
```

```
In [5]: # View the column names
data.columns
```

```
Out[5]: Index(['Time', 'V1', 'V2', 'V3', 'V4', 'V5', 'V6', 'V7', 'V8', 'V9', 'V10',
              'V11', 'V12', 'V13', 'V14', 'V15', 'V16', 'V17', 'V18', 'V19', 'V20',
              'V21', 'V22', 'V23', 'V24', 'V25', 'V26', 'V27', 'V28', 'Amount',
              'Class'],
              dtype='object')
```

```
In [6]: # View the first 10 values
data.head(10)
```

Out[6]:

	Time	V1	V2	V3	V4	V5	V6	V7	V8	
0	0.0	-1.359807	-0.072781	2.536347	1.378155	-0.338321	0.462388	0.239599	0.098698	
1	0.0	1.191857	0.266151	0.166480	0.448154	0.060018	-0.082361	-0.078803	0.085102	-
2	1.0	-1.358354	-1.340163	1.773209	0.379780	-0.503198	1.800499	0.791461	0.247676	-
3	1.0	-0.966272	-0.185226	1.792993	-0.863291	-0.010309	1.247203	0.237609	0.377436	-
4	2.0	-1.158233	0.877737	1.548718	0.403034	-0.407193	0.095921	0.592941	-0.270533	
5	2.0	-0.425966	0.960523	1.141109	-0.168252	0.420987	-0.029728	0.476201	0.260314	-
6	4.0	1.229658	0.141004	0.045371	1.202613	0.191881	0.272708	-0.005159	0.081213	
7	7.0	-0.644269	1.417964	1.074380	-0.492199	0.948934	0.428118	1.120631	-3.807864	
8	7.0	-0.894286	0.286157	-0.113192	-0.271526	2.669599	3.721818	0.370145	0.851084	-
9	9.0	-0.338262	1.119593	1.044367	-0.222187	0.499361	-0.246761	0.651583	0.069539	-

10 rows × 31 columns

```
In [7]: # View the last 10 values
data.tail(10)
```

Out[7]:

	Time	V1	V2	V3	V4	V5	V6	V7	
284797	172782.0	-0.241923	0.712247	0.399806	-0.463406	0.244531	-1.343668	0.929369	-(
284798	172782.0	0.219529	0.881246	-0.635891	0.960928	-0.152971	-1.014307	0.427126	(
284799	172783.0	-1.775135	-0.004235	1.189786	0.331096	1.196063	5.519980	-1.518185	;
284800	172784.0	2.039560	-0.175233	-1.196825	0.234580	-0.008713	-0.726571	0.017050	-(
284801	172785.0	0.120316	0.931005	-0.546012	-0.745097	1.130314	-0.235973	0.812722	(
284802	172786.0	-11.881118	10.071785	-9.834783	-2.066656	-5.364473	-2.606837	-4.918215	;
284803	172787.0	-0.732789	-0.055080	2.035030	-0.738589	0.868229	1.058415	0.024330	(
284804	172788.0	1.919565	-0.301254	-3.249640	-0.557828	2.630515	3.031260	-0.296827	(
284805	172788.0	-0.240440	0.530483	0.702510	0.689799	-0.377961	0.623708	-0.686180	(
284806	172792.0	-0.533413	-0.189733	0.703337	-0.506271	-0.012546	-0.649617	1.577006	-(

10 rows × 31 columns

```
In [8]: # Check if any null values in the data  
data.isnull().any()
```

```
Out[8]: Time      False  
V1          False  
V2          False  
V3          False  
V4          False  
V5          False  
V6          False  
V7          False  
V8          False  
V9          False  
V10         False  
V11         False  
V12         False  
V13         False  
V14         False  
V15         False  
V16         False  
V17         False  
V18         False  
V19         False  
V20         False  
V21         False  
V22         False  
V23         False  
V24         False  
V25         False  
V26         False  
V27         False  
V28         False  
Amount      False  
Class       False  
dtype: bool
```

```
In [9]: data.isnull().sum()
```

```
Out[9]: Time      0
        V1        0
        V2        0
        V3        0
        V4        0
        V5        0
        V6        0
        V7        0
        V8        0
        V9        0
        V10       0
        V11       0
        V12       0
        V13       0
        V14       0
        V15       0
        V16       0
        V17       0
        V18       0
        V19       0
        V20       0
        V21       0
        V22       0
        V23       0
        V24       0
        V25       0
        V26       0
        V27       0
        V28       0
        Amount    0
        Class     0
        dtype: int64
```

No null values in the given dataset

Exploratory Data Analysis(EDA)

```
In [10]: data.describe().transpose()
```

Out[10]:

	count	mean	std	min	25%	50%	
Time	284807.0	9.481386e+04	47488.145955	0.000000	54201.500000	84692.000000	139320.0
V1	284807.0	3.919560e-15	1.958696	-56.407510	-0.920373	0.018109	1.0
V2	284807.0	5.688174e-16	1.651309	-72.715728	-0.598550	0.065486	0.0
V3	284807.0	-8.769071e-15	1.516255	-48.325589	-0.890365	0.179846	1.0
V4	284807.0	2.782312e-15	1.415869	-5.683171	-0.848640	-0.019847	0.0
V5	284807.0	-1.552563e-15	1.380247	-113.743307	-0.691597	-0.054336	0.0
V6	284807.0	2.010663e-15	1.332271	-26.160506	-0.768296	-0.274187	0.0
V7	284807.0	-1.694249e-15	1.237094	-43.557242	-0.554076	0.040103	0.0
V8	284807.0	-1.927028e-16	1.194353	-73.216718	-0.208630	0.022358	0.0
V9	284807.0	-3.137024e-15	1.098632	-13.434066	-0.643098	-0.051429	0.0
V10	284807.0	1.768627e-15	1.088850	-24.588262	-0.535426	-0.092917	0.0
V11	284807.0	9.170318e-16	1.020713	-4.797473	-0.762494	-0.032757	0.0
V12	284807.0	-1.810658e-15	0.999201	-18.683715	-0.405571	0.140033	0.0
V13	284807.0	1.693438e-15	0.995274	-5.791881	-0.648539	-0.013568	0.0
V14	284807.0	1.479045e-15	0.958596	-19.214325	-0.425574	0.050601	0.0
V15	284807.0	3.482336e-15	0.915316	-4.498945	-0.582884	0.048072	0.0
V16	284807.0	1.392007e-15	0.876253	-14.129855	-0.468037	0.066413	0.0
V17	284807.0	-7.528491e-16	0.849337	-25.162799	-0.483748	-0.065676	0.0
V18	284807.0	4.328772e-16	0.838176	-9.498746	-0.498850	-0.003636	0.0
V19	284807.0	9.049732e-16	0.814041	-7.213527	-0.456299	0.003735	0.0
V20	284807.0	5.085503e-16	0.770925	-54.497720	-0.211721	-0.062481	0.0
V21	284807.0	1.537294e-16	0.734524	-34.830382	-0.228395	-0.029450	0.0
V22	284807.0	7.959909e-16	0.725702	-10.933144	-0.542350	0.006782	0.0
V23	284807.0	5.367590e-16	0.624460	-44.807735	-0.161846	-0.011193	0.0
V24	284807.0	4.458112e-15	0.605647	-2.836627	-0.354586	0.040976	0.0
V25	284807.0	1.453003e-15	0.521278	-10.295397	-0.317145	0.016594	0.0
V26	284807.0	1.699104e-15	0.482227	-2.604551	-0.326984	-0.052139	0.0
V27	284807.0	-3.660161e-16	0.403632	-22.565679	-0.070840	0.001342	0.0
V28	284807.0	-1.206049e-16	0.330083	-15.430084	-0.052960	0.011244	0.0
Amount	284807.0	8.834962e+01	250.120109	0.000000	5.600000	22.000000	77.0
Class	284807.0	1.727486e-03	0.041527	0.000000	0.000000	0.000000	0.0

```
In [11]: # Find how many 0,1 in Out put data (Class)
data['Class'].value_counts()
```

```
Out[11]: 0    284315
         1      492
         Name: Class, dtype: int64
```

For predicting the model ,we have to take 0 as non_Fraud and 1 as Fraud happening.

Feature Scaling

```
In [12]: from sklearn.preprocessing import StandardScaler
data['normalizedAmount'] = StandardScaler().fit_transform(data['Amount'].values.reshape(-1,1))
data = data.drop(['Amount'],axis=1)
```

```
In [13]: data = data.drop(['Time'],axis=1)
data.head()
```

```
Out[13]:
```


	V1	V2	V3	V4	V5	V6	V7	V8	V
0	-1.359807	-0.072781	2.536347	1.378155	-0.338321	0.462388	0.239599	0.098698	0.36378
1	1.191857	0.266151	0.166480	0.448154	0.060018	-0.082361	-0.078803	0.085102	-0.25542
2	-1.358354	-1.340163	1.773209	0.379780	-0.503198	1.800499	0.791461	0.247676	-1.51465
3	-0.966272	-0.185226	1.792993	-0.863291	-0.010309	1.247203	0.237609	0.377436	-1.38702
4	-1.158233	0.877737	1.548718	0.403034	-0.407193	0.095921	0.592941	-0.270533	0.81773

5 rows × 30 columns

Data Visualizations


1. Univariate Analysis


```
In [14]: fig,ax = plt.subplots(figsize=(15,5))
ax = sns.countplot(data['Class'])
plt.show()
```



Distribution of the Class


```
In [15]: # sns.distplot(data['Class'])
sns.distplot(data['Class'],rug=True)
plt.show()
```




Correlation

```
In [16]: data.corrwith(data.Class).plot.bar(figsize = (20, 10), title = "Correlation with class", fontsize = 15, rot = 45, grid = True)
```

```
Out[16]: <matplotlib.axes._subplots.AxesSubplot at 0x12f5328b448>
```



```
In [17]: plt.figure(figsize=(14,10))
sns.heatmap(data.corr(),annot=True,cmap='hsv',fmt='.3f',linewidths=2)
plt.show()
```



```
In [18]: # Compute the correlation matrix
corr = data.corr()
corr.head()
```

```
Out[18]:
```

	V1	V2	V3	V4	V5	V6
V1	1.000000e+00	4.697350e-17	-1.424390e-15	1.755316e-17	6.391162e-17	2.398071e-16
V2	4.697350e-17	1.000000e+00	2.512175e-16	-1.126388e-16	-2.039868e-16	5.024680e-16
V3	-1.424390e-15	2.512175e-16	1.000000e+00	-3.416910e-16	-1.436514e-15	1.431581e-15
V4	1.755316e-17	-1.126388e-16	-3.416910e-16	1.000000e+00	-1.940929e-15	-2.712659e-16
V5	6.391162e-17	-2.039868e-16	-1.436514e-15	-1.940929e-15	1.000000e+00	7.926364e-16

5 rows × 30 columns


```
In [19]: # Generate a mask for the upper triangle
mask = np.zeros_like(corr, dtype=np.bool)
mask[np.triu_indices_from(mask)] = True

# Set up the matplotlib figure
f, ax = plt.subplots(figsize=(18, 15))

# Generate a custom diverging colormap
cmap = sns.diverging_palette(220, 10, as_cmap=True)

# Draw the heatmap with the mask and correct aspect ratio
sns.heatmap(corr, mask=mask, cmap=cmap, vmax=.3, center=0, square=True, linewidth
ths=.5, cbar_kws={"shrink": .5})
```

Out[19]: <matplotlib.axes._subplots.AxesSubplot at 0x12f67aa0e08>



Normalize the data

```
In [20]: from sklearn.preprocessing import scale
scale_data=scale(data)
scale_data
```

```
Out[20]: array([[ -0.69424232, -0.04407492,  1.6727735 , ..., -0.06378115,
        -0.04159898,  0.24496426],
       [ 0.60849633,  0.16117592,  0.1097971 , ...,  0.04460752,
        -0.04159898, -0.34247454],
       [-0.69350046, -0.81157783,  1.16946849, ..., -0.18102083,
        -0.04159898,  1.16068593],
       ...,
       [ 0.98002374, -0.18243372, -2.14320514, ..., -0.0804672 ,
        -0.04159898, -0.0818393 ],
       [-0.12275539,  0.32125034,  0.46332013, ...,  0.31668678,
        -0.04159898, -0.31324853],
       [-0.27233093, -0.11489898,  0.46386564, ...,  0.04134999,
        -0.04159898,  0.51435531]])
```

```
In [21]: np.exp(scale_data)
```

```
Out[21]: array([[0.49945273, 0.95688226, 5.32692154, ..., 0.9382103 , 0.95925439,
        1.27757566],
       [1.83766607, 1.17489164, 1.1160516 , ..., 1.04561739, 0.95925439,
        0.7100112 ],
       [0.49982339, 0.44415671, 3.22028058, ..., 0.83441798, 0.95925439,
        3.19212208],
       ...,
       [2.66451949, 0.83323987, 0.11727835, ..., 0.92268517, 0.95925439,
        0.92142002],
       [0.88447999, 1.37885072, 1.58934205, ..., 1.37257258, 0.95925439,
        0.7310682 ],
       [0.76160218, 0.89145619, 1.5902093 , ..., 1.0422168 , 0.95925439,
        1.67255987]])
```

Divide the dataset into input and output

```
In [22]: X = data.iloc[:,data.columns!= 'Class']
Y = data.iloc[:,data.columns== 'Class']
```

```
In [23]: # View the input data
X.head()
```

```
Out[23]:
```

	V1	V2	V3	V4	V5	V6	V7	V8	V
0	-1.359807	-0.072781	2.536347	1.378155	-0.338321	0.462388	0.239599	0.098698	0.36378
1	1.191857	0.266151	0.166480	0.448154	0.060018	-0.082361	-0.078803	0.085102	-0.25542
2	-1.358354	-1.340163	1.773209	0.379780	-0.503198	1.800499	0.791461	0.247676	-1.51465
3	-0.966272	-0.185226	1.792993	-0.863291	-0.010309	1.247203	0.237609	0.377436	-1.38702
4	-1.158233	0.877737	1.548718	0.403034	-0.407193	0.095921	0.592941	-0.270533	0.81773

5 rows × 29 columns

```
In [24]: # View the out put Data
Y.head()
```

```
Out[24]:
```

	Class
0	0
1	0
2	0
3	0
4	0

Split the data as train and test

```
In [25]: from sklearn.model_selection import train_test_split
x_train,x_test,y_train,y_test = train_test_split(X,Y,train_size=0.80,random_state=0)
```

Build the Models

1. Decision Tree Classifier

```
In [26]: from sklearn.tree import DecisionTreeClassifier
classifier = DecisionTreeClassifier(random_state = 0,criterion = 'gini', splitter='best', min_samples_leaf=1, min_samples_split=2)
```

```
In [27]: # Fit the model
classifier.fit(x_train,y_train)
```

```
Out[27]: DecisionTreeClassifier(class_weight=None, criterion='gini', max_depth=None,
                                max_features=None, max_leaf_nodes=None,
                                min_impurity_decrease=0.0, min_impurity_split=None,
                                min_samples_leaf=1, min_samples_split=2,
                                min_weight_fraction_leaf=0.0, presort=False,
                                random_state=0, splitter='best')
```

```
In [28]: # Predict the model
DT_pred = classifier.predict(x_test)
DT_pred
```

```
Out[28]: array([0, 0, 0, ..., 0, 0, 0], dtype=int64)
```

```
In [29]: # Confusion Matrix
from sklearn.metrics import confusion_matrix,accuracy_score,f1_score,precision
_score,recall_score
print(confusion_matrix(DT_pred,y_test))
```

```
[[56839   26]
 [   22   75]]
```

```
In [30]: # Accuracy Score
DT_acc = accuracy_score(DT_pred,y_test)
DT_acc
```

```
Out[30]: 0.9991573329588147
```

```
In [31]: # Precision ,Recall,F1_score
DT_Prec = precision_score(DT_pred,y_test)
DT_rec = recall_score(DT_pred,y_test)
DT_f1 = f1_score(DT_pred,y_test)
```

```
In [32]: # Store the results
results = pd.DataFrame([[ 'Decision tree', DT_acc, DT_Prec, DT_rec, DT_f1]],col
umns = [ 'Model', 'Accuracy', 'Precision', 'Recall', 'F1 Score'])
```

```
In [33]: # view the results
results
```

```
Out[33]:
```

	Model	Accuracy	Precision	Recall	F1 Score
0	Decision tree	0.999157	0.742574	0.773196	0.757576

2. Random Forest Classifier

```
In [34]: from sklearn.ensemble import RandomForestClassifier
classifier1 = RandomForestClassifier(random_state = 0, n_estimators = 100, criterion = 'entropy')
classifier1
```

```
Out[34]: RandomForestClassifier(bootstrap=True, class_weight=None, criterion='entropy',
                                max_depth=None, max_features='auto', max_leaf_nodes=None,
                                min_impurity_decrease=0.0, min_impurity_split=None,
                                min_samples_leaf=1, min_samples_split=2,
                                min_weight_fraction_leaf=0.0, n_estimators=100,
                                n_jobs=None, oob_score=False, random_state=0, verbose=0,
                                warm_start=False)
```

```
In [35]: # Fit the model
classifier1.fit(x_train,y_train)
```

C:\Users\sundara.rao.ext\Anaconda\lib\site-packages\ipykernel_launcher.py:2:
DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n_samples,), for example using ravel().

```
Out[35]: RandomForestClassifier(bootstrap=True, class_weight=None, criterion='entropy',
                                max_depth=None, max_features='auto', max_leaf_nodes=None,
                                min_impurity_decrease=0.0, min_impurity_split=None,
                                min_samples_leaf=1, min_samples_split=2,
                                min_weight_fraction_leaf=0.0, n_estimators=100,
                                n_jobs=None, oob_score=False, random_state=0, verbose=0,
                                warm_start=False)
```

```
In [36]: # predict the model
RF_pred = classifier1.predict(x_test)
RF_pred
```

```
Out[36]: array([0, 0, 0, ..., 0, 0, 0], dtype=int64)
```

```
In [37]: # confusion_matrix
print(confusion_matrix(RF_pred,y_test))
```

```
[[56855   22]
 [    6   79]]
```

```
In [38]: # Accuracy , precision,recall,f1_score
RF_acc = accuracy_score(RF_pred,y_test)
RF_prec = precision_score(RF_pred,y_test)
RF_rec = recall_score(RF_pred,y_test)
RF_f1= f1_score(RF_pred,y_test)
```



```
In [39]: # Store the results
results1 = pd.DataFrame([[ 'Random Forest (n=100)', RF_acc, RF_prec, RF_rec, RF_f1]],columns = [ 'Model', 'Accuracy', 'Precision', 'Recall', 'F1 Score'])
```

```
In [40]: # view the results
results1
```

```
Out[40]:
```

	Model	Accuracy	Precision	Recall	F1 Score
0	Random Forest (n=100)	0.999508	0.782178	0.929412	0.849462

3.Naive Bayes Classifier

```
In [41]: from sklearn.naive_bayes import GaussianNB
classifier2 = GaussianNB()
```

```
In [42]: # Fit the model
classifier2.fit(x_train,y_train)
```

C:\Users\sundara.rao.ext\Anaconda\lib\site-packages\sklearn\utils\validation.py:724: DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n_samples,), for example using ravel().

```
y = column_or_1d(y, warn=True)
```

```
Out[42]: GaussianNB(priors=None, var_smoothing=1e-09)
```

```
In [43]: #Predict the model
NBC_pred = classifier2.predict(x_test)
NBC_pred
```

```
Out[43]: array([0, 0, 0, ..., 0, 0, 0], dtype=int64)
```

```
In [44]: # Confusion Matrix
print(confusion_matrix(NBC_pred,y_test))
```

```
[[55642   15]
 [ 1219   86]]
```

```
In [45]: # Accuracy ,Precision,Recall,f1_core
NBC_acc = accuracy_score(NBC_pred,y_test)
NBC_prec = precision_score(NBC_pred,y_test)
NBC_rec = recall_score(NBC_pred,y_test)
NBC_f1= f1_score(NBC_pred,y_test)
```

```
In [46]: # Store the results
results2 = pd.DataFrame([[ 'Naive Bayes', NBC_acc, NBC_prec, NBC_rec, NBC_f1]],
columns = [ 'Model', 'Accuracy', 'Precision', 'Recall', 'F1 Score'])
```

```
In [47]: # view the result
results2
```

Out[47]:

	Model	Accuracy	Precision	Recall	F1 Score
0	Naive Bayes	0.978336	0.851485	0.0659	0.122333

K-Nearest Neighbor

```
In [48]: from sklearn.neighbors import KNeighborsClassifier
classifier3 = KNeighborsClassifier(n_neighbors=5)
```

```
In [49]: # Fit the model
classifier3.fit(x_train,y_train)
```

C:\Users\sundara.rao.ext\Anaconda\lib\site-packages\ipykernel_launcher.py:2:
DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n_samples,), for example using ravel().

Out[49]: KNeighborsClassifier(algorithm='auto', leaf_size=30, metric='minkowski',
metric_params=None, n_jobs=None, n_neighbors=5, p=2,
weights='uniform')

```
In [50]: # Predict the model
KNN_pred = classifier3.predict(x_test)
KNN_pred
```

Out[50]: array([0, 0, 0, ..., 0, 0, 0], dtype=int64)

```
In [51]: # Confusion Matrix
print(confusion_matrix(KNN_pred,y_test))
```

```
[[56854  20]
 [    7   81]]
```

```
In [52]: # Accuracy ,Precision,Recall,f1_core
KNN_acc = accuracy_score(KNN_pred,y_test)
KNN_prec = precision_score(KNN_pred,y_test)
KNN_rec = recall_score(KNN_pred,y_test)
KNN_f1= f1_score(KNN_pred,y_test)
```

```
In [53]: # Store the results
results3 = pd.DataFrame([['K-Nearest Neighbor', KNN_acc, KNN_prec, KNN_rec, KNN_f1]],columns = ['Model', 'Accuracy', 'Precision', 'Recall', 'F1 Score'])
```

```
In [54]: # view the result
results3
```

Out[54]:

	Model	Accuracy	Precision	Recall	F1 Score
0	K-Nearest Neighbor	0.999526	0.80198	0.920455	0.857143

5. Support Vector Machine (SVM)

```
In [55]: from sklearn.svm import SVC
classifier4 = SVC(kernel='poly', random_state =0)
```

```
In [56]: # fit the model
classifier4.fit(x_train,y_train)
```

C:\Users\sundara.rao.ext\Anaconda\lib\site-packages\sklearn\utils\validation.py:724: DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n_samples,), for example using ravel().

y = column_or_1d(y, warn=True)

C:\Users\sundara.rao.ext\Anaconda\lib\site-packages\sklearn\svm\base.py:193: FutureWarning: The default value of gamma will change from 'auto' to 'scale' in version 0.22 to account better for unscaled features. Set gamma explicitly to 'auto' or 'scale' to avoid this warning.

"avoid this warning.", FutureWarning)

Out[56]: SVC(C=1.0, cache_size=200, class_weight=None, coef0=0.0, decision_function_shape='ovr', degree=3, gamma='auto_deprecated', kernel='poly', max_iter=-1, probability=False, random_state=0, shrinking=True, tol=0.001, verbose=False)

```
In [57]: # Predict the model
SVM_pred = classifier4.predict(x_test)
```

```
In [58]: # Confusion Matrix
print(confusion_matrix(SVM_pred,y_test))
```

```
[[56849  22]
 [  12  79]]
```

```
In [59]: # Accuracy ,Precision,Recall,f1_core
SVM_acc = accuracy_score(SVM_pred,y_test)
SVM_prec = precision_score(SVM_pred,y_test)
SVM_rec = recall_score(SVM_pred,y_test)
SVM_f1= f1_score(SVM_pred,y_test)
```

```
In [60]: # Store the results
results4 = pd.DataFrame([['Support Vector Machine', SVM_acc, SVM_prec, SVM_rec, SVM_f1]],columns = ['Model', 'Accuracy', 'Precision', 'Recall', 'F1 Score'])
```

In [61]: results4

Out[61]:

	Model	Accuracy	Precision	Recall	F1 Score
0	Support Vector Machine	0.999403	0.782178	0.868132	0.822917

Artificial Neural Networks

```
In [62]: # Importing the Keras Libraries and packages
import keras
from keras.models import Sequential
from keras.layers import Dense

# Initialising the ANN
classifier = Sequential()

# Adding the input layer and the first hidden layer
classifier.add(Dense(units = 15 , kernel_initializer = 'uniform', activation =
'relu', input_dim = 29))

# Adding the second hidden layer
classifier.add(Dense(units = 15, kernel_initializer = 'uniform', activation =
'relu'))

# Adding the output layer
classifier.add(Dense(units = 1, kernel_initializer = 'uniform', activation =
'sigmoid'))

# Compiling the ANN
classifier.compile(optimizer = 'adam', loss = 'binary_crossentropy', metrics =
['accuracy'])
```

Using TensorFlow backend.

```
In [64]: # Fitting the ANN to the Training set  
classifier.fit(x_train, y_train, batch_size = 32, epochs = 100)
```

Epoch 1/100
227845/227845 [=====] - 14s 63us/step - loss: 0.0132
- accuracy: 0.9986

Epoch 2/100
227845/227845 [=====] - 14s 61us/step - loss: 0.0030
- accuracy: 0.9994

Epoch 3/100
227845/227845 [=====] - 14s 63us/step - loss: 0.0029
- accuracy: 0.9994

Epoch 4/100
227845/227845 [=====] - 14s 63us/step - loss: 0.0027
- accuracy: 0.9994

Epoch 5/100
227845/227845 [=====] - 14s 63us/step - loss: 0.0027
- accuracy: 0.9994

Epoch 6/100
227845/227845 [=====] - 17s 74us/step - loss: 0.0026
- accuracy: 0.9994

Epoch 7/100
227845/227845 [=====] - 15s 64us/step - loss: 0.0027
- accuracy: 0.9994

Epoch 8/100
227845/227845 [=====] - 13s 58us/step - loss: 0.0025
- accuracy: 0.9994

Epoch 9/100
227845/227845 [=====] - 19s 84us/step - loss: 0.0025
- accuracy: 0.9994

Epoch 10/100
227845/227845 [=====] - 18s 77us/step - loss: 0.0025
- accuracy: 0.9995

Epoch 11/100
227845/227845 [=====] - 16s 70us/step - loss: 0.0025
- accuracy: 0.9994

Epoch 12/100
227845/227845 [=====] - 16s 69us/step - loss: 0.0023
- accuracy: 0.9995

Epoch 13/100
227845/227845 [=====] - 16s 71us/step - loss: 0.0024
- accuracy: 0.9995

Epoch 14/100
227845/227845 [=====] - 16s 71us/step - loss: 0.0023
- accuracy: 0.9995

Epoch 15/100
227845/227845 [=====] - 17s 73us/step - loss: 0.0023
- accuracy: 0.9995

Epoch 16/100
227845/227845 [=====] - 19s 83us/step - loss: 0.0022
- accuracy: 0.9995

Epoch 17/100
227845/227845 [=====] - 17s 75us/step - loss: 0.0023
- accuracy: 0.9995

Epoch 18/100
227845/227845 [=====] - 16s 72us/step - loss: 0.0023
- accuracy: 0.9995

Epoch 19/100
227845/227845 [=====] - 22s 95us/step - loss: 0.0022
- accuracy: 0.9995

Epoch 20/100
227845/227845 [=====] - 23s 100us/step - loss: 0.002
2 - accuracy: 0.9995

Epoch 21/100
227845/227845 [=====] - 16s 70us/step - loss: 0.0021
- accuracy: 0.9995

Epoch 22/100
227845/227845 [=====] - 15s 64us/step - loss: 0.0021
- accuracy: 0.9995

Epoch 23/100
227845/227845 [=====] - 13s 56us/step - loss: 0.0020
- accuracy: 0.9995

Epoch 24/100
227845/227845 [=====] - 14s 59us/step - loss: 0.0022
- accuracy: 0.9995

Epoch 25/100
227845/227845 [=====] - 13s 57us/step - loss: 0.0020
- accuracy: 0.9996

Epoch 26/100
227845/227845 [=====] - 13s 58us/step - loss: 0.0020
- accuracy: 0.9995

Epoch 27/100
227845/227845 [=====] - 13s 59us/step - loss: 0.0020
- accuracy: 0.9995

Epoch 28/100
227845/227845 [=====] - 13s 58us/step - loss: 0.0020
- accuracy: 0.9995

Epoch 29/100
227845/227845 [=====] - 13s 57us/step - loss: 0.0019
- accuracy: 0.9995

Epoch 30/100
227845/227845 [=====] - 14s 62us/step - loss: 0.0019
- accuracy: 0.9995

Epoch 38/100
227845/227845 [=====] - 14s 62us/step - loss: 0.0018
- accuracy: 0.9995

Epoch 39/100
227845/227845 [=====] - 14s 62us/step - loss: 0.0018
- accuracy: 0.9995

Epoch 40/100
227845/227845 [=====] - 15s 65us/step - loss: 0.0017
- accuracy: 0.9996

Epoch 41/100
227845/227845 [=====] - 14s 62us/step - loss: 0.0018
- accuracy: 0.9996

Epoch 42/100
227845/227845 [=====] - 14s 60us/step - loss: 0.0017
- accuracy: 0.9996

Epoch 43/100
227845/227845 [=====] - 14s 60us/step - loss: 0.0017
- accuracy: 0.9995

Epoch 44/100
227845/227845 [=====] - 14s 60us/step - loss: 0.0017
- accuracy: 0.9996

Epoch 45/100
227845/227845 [=====] - 13s 56us/step - loss: 0.0017
- accuracy: 0.9996

Epoch 46/100
227845/227845 [=====] - 13s 56us/step - loss: 0.0018
- accuracy: 0.9995

Epoch 47/100
227845/227845 [=====] - 13s 57us/step - loss: 0.0016
- accuracy: 0.9996

Epoch 48/100
227845/227845 [=====] - 13s 58us/step - loss: 0.0016
- accuracy: 0.9995

Epoch 49/100
227845/227845 [=====] - 13s 57us/step - loss: 0.0017
- accuracy: 0.9996

Epoch 50/100
227845/227845 [=====] - 13s 59us/step - loss: 0.0016
- accuracy: 0.9996

Epoch 51/100
227845/227845 [=====] - 16s 70us/step - loss: 0.0016
- accuracy: 0.9996

Epoch 52/100
227845/227845 [=====] - 18s 80us/step - loss: 0.0016
- accuracy: 0.9996

Epoch 53/100
227845/227845 [=====] - 16s 68us/step - loss: 0.0016
- accuracy: 0.9996

Epoch 54/100
227845/227845 [=====] - 15s 68us/step - loss: 0.0016
- accuracy: 0.9995

Epoch 55/100
227845/227845 [=====] - 18s 78us/step - loss: 0.0016
- accuracy: 0.9996

Epoch 56/100
227845/227845 [=====] - 17s 73us/step - loss: 0.0016
- accuracy: 0.9996

Epoch 57/100
227845/227845 [=====] - 17s 74us/step - loss: 0.0015
- accuracy: 0.9996

Epoch 58/100
227845/227845 [=====] - 16s 71us/step - loss: 0.0016
- accuracy: 0.9996

Epoch 59/100
227845/227845 [=====] - 17s 75us/step - loss: 0.0016
- accuracy: 0.9996

Epoch 60/100
227845/227845 [=====] - 17s 73us/step - loss: 0.0015
- accuracy: 0.9996

Epoch 61/100
227845/227845 [=====] - 16s 72us/step - loss: 0.0016
- accuracy: 0.9996

Epoch 62/100
227845/227845 [=====] - 16s 72us/step - loss: 0.0015
- accuracy: 0.9996

Epoch 63/100
227845/227845 [=====] - 13s 59us/step - loss: 0.0015
- accuracy: 0.9996

Epoch 64/100
227845/227845 [=====] - 15s 66us/step - loss: 0.0014
- accuracy: 0.9996

Epoch 65/100
227845/227845 [=====] - 16s 69us/step - loss: 0.0015
- accuracy: 0.9996

Epoch 66/100
227845/227845 [=====] - 15s 65us/step - loss: 0.0015
- accuracy: 0.9996

Epoch 67/100
227845/227845 [=====] - 15s 65us/step - loss: 0.0016
- accuracy: 0.9996

Epoch 68/100
227845/227845 [=====] - 13s 56us/step - loss: 0.0015
- accuracy: 0.9996

Epoch 69/100
227845/227845 [=====] - 13s 59us/step - loss: 0.0015
- accuracy: 0.9996

Epoch 70/100
227845/227845 [=====] - 15s 66us/step - loss: 0.0015
- accuracy: 0.9996

Epoch 71/100
227845/227845 [=====] - 13s 56us/step - loss: 0.0014
- accuracy: 0.9996

Epoch 72/100
227845/227845 [=====] - 13s 56us/step - loss: 0.0015
- accuracy: 0.9996

Epoch 73/100
227845/227845 [=====] - 12s 53us/step - loss: 0.0015
- accuracy: 0.9996

Epoch 74/100
227845/227845 [=====] - 12s 51us/step - loss: 0.0015
- accuracy: 0.9996

Epoch 75/100
227845/227845 [=====] - 12s 53us/step - loss: 0.0015
- accuracy: 0.9996

Epoch 76/100
227845/227845 [=====] - 12s 54us/step - loss: 0.0014
- accuracy: 0.9996

Epoch 77/100
227845/227845 [=====] - 12s 54us/step - loss: 0.0014
- accuracy: 0.9996

Epoch 78/100
227845/227845 [=====] - 16s 68us/step - loss: 0.0014
- accuracy: 0.9996

Epoch 79/100
227845/227845 [=====] - 17s 74us/step - loss: 0.0014
- accuracy: 0.9996

Epoch 80/100
227845/227845 [=====] - 17s 76us/step - loss: 0.0014
- accuracy: 0.9996

Epoch 81/100
227845/227845 [=====] - 17s 75us/step - loss: 0.0014
- accuracy: 0.9996

Epoch 82/100
227845/227845 [=====] - 16s 68us/step - loss: 0.0014
- accuracy: 0.9996

Epoch 83/100
227845/227845 [=====] - 16s 72us/step - loss: 0.0014
- accuracy: 0.9996

```
Epoch 84/100
227845/227845 [=====] - 16s 69us/step - loss: 0.0014
- accuracy: 0.9996
Epoch 85/100
227845/227845 [=====] - 16s 70us/step - loss: 0.0014
- accuracy: 0.9996
Epoch 86/100
227845/227845 [=====] - 16s 72us/step - loss: 0.0014
- accuracy: 0.9996
Epoch 87/100
227845/227845 [=====] - 17s 72us/step - loss: 0.0014
- accuracy: 0.9996
Epoch 88/100
227845/227845 [=====] - 16s 70us/step - loss: 0.0014
- accuracy: 0.9996
Epoch 89/100
227845/227845 [=====] - 16s 70us/step - loss: 0.0013
- accuracy: 0.9996
Epoch 90/100
227845/227845 [=====] - 16s 71us/step - loss: 0.0013
- accuracy: 0.9996
Epoch 91/100
227845/227845 [=====] - 16s 70us/step - loss: 0.0016
- accuracy: 0.9995
Epoch 92/100
227845/227845 [=====] - 17s 74us/step - loss: 0.0014
- accuracy: 0.9996
Epoch 93/100
227845/227845 [=====] - 13s 55us/step - loss: 0.0014
- accuracy: 0.9996
Epoch 94/100
227845/227845 [=====] - 13s 57us/step - loss: 0.0015
- accuracy: 0.9996
Epoch 95/100
227845/227845 [=====] - 13s 58us/step - loss: 0.0014
- accuracy: 0.9996
Epoch 96/100
227845/227845 [=====] - 13s 58us/step - loss: 0.0013
- accuracy: 0.9996
Epoch 97/100
227845/227845 [=====] - 13s 58us/step - loss: 0.0013
- accuracy: 0.9996
Epoch 98/100
227845/227845 [=====] - 13s 57us/step - loss: 0.0013
- accuracy: 0.9996
Epoch 99/100
227845/227845 [=====] - 12s 54us/step - loss: 0.0013
- accuracy: 0.9996
Epoch 100/100
227845/227845 [=====] - 12s 52us/step - loss: 0.0013
- accuracy: 0.9996
```

Out[64]: <keras.callbacks.callbacks.History at 0x12f757b1e48>

```
In [66]: # Predicting the Test set results
y_pred = classifier.predict(x_test)
y_pred = (y_pred > 0.5)
```

```
In [67]: score = classifier.evaluate(x_test, y_test)
score
```

56962/56962 [=====] - 2s 32us/step

```
Out[67]: [0.0049244485128219564, 0.9993504285812378]
```

```
In [69]: # Making the Confusion Matrix
cm = confusion_matrix(y_test, y_pred)
cm
```


```
Out[69]: array([[56843,   18],
               [   19,   82]], dtype=int64)
```

```
In [70]: #Let's see how our model performed
from sklearn.metrics import classification_report
print(classification_report(y_test, y_pred))
```

	precision	recall	f1-score	support
0	1.00	1.00	1.00	56861
1	0.82	0.81	0.82	101
accuracy			1.00	56962
macro avg	0.91	0.91	0.91	56962
weighted avg	1.00	1.00	1.00	56962

```
In [72]: ## EXTRA: Confusion Matrix
cm = confusion_matrix(y_test, y_pred) # rows = truth, cols = prediction
df_cm = pd.DataFrame(cm, index = (0, 1), columns = (0, 1))
plt.figure(figsize = (10,7))
sns.set(font_scale=1.4)
sns.heatmap(df_cm, annot=True, fmt='g')
print("Test Data Accuracy: %0.4f" % accuracy_score(y_test, y_pred))
```

Test Data Accuracy: 0.9994



In []: